

Passing Is Not Shipping

Correctness, Semantic Equivalence and ReDoS Safety in
LLM-Generated Regular Expressions, and the Limits of Measuring Them

Foothills Labs
info@foothills-labs.com

August 2026

Abstract

We evaluate eleven contemporary language models on natural-language-to-regex generation over 450 tasks with three samples each (14,850 requests, \$10.95), scoring every answer on three axes: whether it satisfies the task’s tests, whether it denotes the same language as a human reference, and whether it is vulnerable to regular-expression denial of service (ReDoS).

Two empirical results are robust. First, *passing is roughly twice as easy as shipping*: models pass 38.0–47.5% of tasks but produce a correct, safe, semantically-verified pattern on only 17.1–24.8%, a gap that holds uniformly across a 98× range in price. Second, and contrary to the expected framing, *models are safer than the humans whose answers define the benchmark*: 13.6% of the corpus’s human-written reference patterns are ReDoS-vulnerable against 9.0% for models pooled, and every individual model is safer than the reference set.

We then audit our own instrument, and report a negative result about benchmark methodology that we believe generalises. Metrics that compare generated output against a human answer key inherit that key’s error rate. In a random sample of disagreements where a model passed every test yet was scored as semantically different, the model was clearly at fault in approximately 15% of cases; the remainder divided between demonstrably incorrect reference patterns (including a reference character class containing a literal `|`) and prompts that never specified the property in dispute. We further show that independent binomial intervals—the default in this literature—are the wrong inferential tool when all systems answer identical items, and that a paired bootstrap resolves nine times as many pairwise comparisons on the same data.

We therefore decline to publish a ranking. We release all raw generations, scoring code and adjudications, and recommend that evaluations separate reference-dependent from reference-independent metrics and assign them different evidentiary weight.

Artifacts. All model outputs, scoring code, per-case adjudications and analysis scripts are available at <https://github.com/foothills-labs/regexleaderboard>. Every number in this paper recomputes offline from committed data with no API access.

1 Introduction

Regular expressions are an unusually good probe for language-model evaluation. The task is compact, the output is executable, correctness is partly decidable, and—uniquely among common code-generation tasks—a syntactically valid, test-passing answer can still be a security defect. That last property is the reason for this study.

The dominant evaluation protocol for natural-language-to-regex generation scores a candidate pattern against a reference using DFA equivalence [Locascio et al., 2016, Kushman and Barzilay,

2013], optionally supplemented by test-based correctness. Neither criterion detects catastrophic backtracking. A pattern that satisfies every provided example and denotes exactly the reference language may still admit inputs on which matching takes time exponential in input length—the ReDoS vulnerability class [Davis et al., 2018]. In deployed systems this is a denial-of-service vector, and it is common enough in human-written code to have motivated ecosystem-scale study.

1.1 Contributions

1. **A three-axis evaluation of eleven current models** (Section 4), reporting test correctness, semantic equivalence and ReDoS safety jointly rather than separately, together with a composite (USABLE) requiring all three. To our knowledge safety has not previously been reported as a first-class axis alongside correctness in this task.
2. **A human baseline for ReDoS** (Section 4.2). We apply the identical safety screen to the corpus’s own reference patterns. Models are uniformly safer than the reference set, which relocates the finding from a claim about model deficiency to a claim about the distribution models were trained on.
3. **A quantified audit of reference-standard error** (Section 5.1). We sample disagreements and adjudicate them, finding that the majority of apparent semantic errors are not model errors. We are not aware of prior work in this task family that measures this.
4. **A methodological correction** (Section 5.2). We show that independent intervals substantially understate resolvable differences when all systems answer identical items, and give the paired alternative.
5. **Full artifact release** including raw generations, which permits any of our scoring decisions to be overturned without re-running inference.

1.2 What this paper does not claim

We do not publish a ranking of the eleven models. Section 5.3 shows that 62% of tasks yield identical outcomes across all eleven systems, and that only nine of fifty-five pairwise comparisons resolve at 95% confidence even under the correct paired test. Reporting an ordinal leaderboard from this data would be unsupported by it.

We also do not claim novelty for DFA-equivalence checking, which is standard in this literature [Locascio et al., 2016], nor for the scoring engine, which is prior work by the present authors released separately [Foothills Labs, 2026] and used here as a pinned dependency.

2 Background and Related Work

NL-to-regex generation. The task was established by Kushman and Barzilay [2013] (KB13) and extended by Locascio et al. [2016] (NL-RX), both of which adopt DFA equivalence against a reference as the correctness criterion. This is the right choice in principle— $[0-9]^+$ and $[0-9][0-9]^*$ denote the same language and differ as strings—but it inherits any defect in the reference set, a dependency this paper quantifies.

ReDoS. Catastrophic backtracking arises when a backtracking matcher can decompose a prefix in exponentially many ways, canonically via a quantifier applied to a quantified group, as in $(a^+)^+$. Davis et al. [2018] document its prevalence at ecosystem scale; Davis et al. [2019] study regex portability and re-use across languages. ICPC 2024 authors — VERIFY [2024] examine ReDoS specifically in LLM-generated patterns and report a skew toward polynomial rather than exponential vulnerability; we compare against that finding in Section 4.2.

Sampling-based metrics. We use the unbiased $PASS@k$ estimator of Chen et al. [2021], with the degeneracy noted in Section 3.4.

Serving-layer variance. Hosted inference routers may serve the same model slug from multiple providers at differing numerical precision. LessWrong [2025] report that 31 of 32 surveyed AI-safety codebases using such a router did not pin the serving provider, making their results a measurement of the router’s routing policy in addition to the model. Section 3.3 describes the discipline we adopt in response.

3 Method

3.1 Task and prompt

Each task supplies a natural-language description; the model returns one regular expression. The prompt is fixed across all models and all tasks:

system: You translate a natural-language description into a single Python **re**-compatible regular expression. Reply with ONLY the pattern in one fenced code block, no explanation.

user: *{task description, verbatim}*

Reply with only the regular expression pattern in a single fenced code block.

We do not tune this prompt. Prompt optimisation would raise all scores and destroy comparability with prior work on the same corpus. Section 6 notes the corresponding threat.

Extraction. We take the final fenced code block, else the whole trimmed reply, then strip at most one layer of host-language quoting (`r'...`, `'...`, `"..."`, backticks, `/.../flags`). Scored literally, `r'\d+$'` is a pattern matching the letter `r`, and would fail for reasons unrelated to regex competence. This affected 7–18 of 1,350 responses per model ($< 1.4\%$ throughout) and changes no conclusion; both normalised and unnormalised scores are released.

3.2 Corpus

We use Re(gEx|DoS)Eval [S2E Lab — VERIFY, 2024], 762 tasks drawn from real user requests, each with a description, positive and negative test strings, and a human-written reference pattern. We evaluate 450 tasks, selected at uniform stride across the corpus rather than as a prefix (the corpus is difficulty-ordered at its head). The subset size was budget-determined.

Match semantics are set per corpus. Re(gEx|DoS)Eval references satisfy their own tests at 100% under search semantics and 94% under full-match; choosing incorrectly would score 46 gold patterns as failures. We verified the loaded configuration by scoring every reference against its own tests before the run (Section 5.1).

3.3 Models and serving discipline

Eleven models spanning frontier and open-weights systems across a 98× price range. Every request pins a single named provider with fallback disabled, so a request is either served by the named endpoint or fails visibly. Substitution failures did occur and are reported as failures rather than silently re-routed (Appendix A).

For open-weights models the pin includes quantisation where the router discloses it: GLM-5.2 and DeepSeek-V4-Flash are served at 4-bit by some providers and 8-bit by others, and we pin 8-bit. One pin went stale within an hour of being set (the provider ceased serving that model), so pins are re-verified immediately before each run.

Crucially, the pin is a request, not a guarantee. We record the provider that *actually served* each response and verify agreement post hoc. All eleven models were served entirely by their pinned endpoint, with no substitutions.

3.4 Metrics

Let $T = (d, P, N, r)$ be a task with description d , positive examples P , negative examples N and reference pattern r . Let p be a candidate.

Correctness.

$$\text{correct}(p, T) = [\forall x \in P : \text{match}(p, x)] \wedge [\forall y \in N : \neg \text{match}(p, y)]$$

evaluated with CPython’s `re` under the corpus’s declared semantics. This is *reference-independent*.

Semantic equivalence. Both p and r are compiled to finite automata and compared, yielding a verdict in $\{\text{EQ}, \text{DIFF}, \text{UNDEC}, \text{UNSUP}\}$. When the verdict is `DIFF` the engine returns a *witness*: a shortest string in the symmetric difference, which renders the judgment falsifiable by inspection.

Equivalence is decidable only on the regular subset. Backreferences render the language non-regular and equivalence formally undecidable, yielding `UNDEC`; between 82 and 133 of 450 tasks per model fall in this class. We therefore report two figures:

$$\text{DFA-EQ} = \Pr[\text{EQ}], \quad \text{DFA-EQ}_{\text{dec}} = \Pr[\text{EQ} \mid \text{verdict} \neq \text{UNDEC}].$$

The first counts undecidable comparisons as failures and is a lower bound that cannot flatter; the second isolates the model from the engine’s reach. Reporting only the latter inflates; reporting only the former charges the model for a theorem. This metric is *reference-dependent*, which Section 5.1 shows to be decisive.

ReDoS safety. $\text{vuln}(p) = [\text{screen}(p) \neq \text{SAFE}]$, where `screen` performs structural analysis for known-vulnerable shapes followed by an empirical pass against attack strings under timeout. The structural pass covers three of the five families catalogued by ICPC 2024 authors — `VERIFY` [2024]; consequently `VULNERABLE` is a *lower bound*, and `SAFE` denotes a screening outcome, not a proof. This metric is *reference-independent*.

Composite.

$$\text{usable}(p, T) = \text{correct}(p, T) \wedge \neg \text{vuln}(p) \wedge [\text{verdict}(p, r) \neq \text{DIFF}]$$

Note the third conjunct excludes only *proven* difference; undecidable and unsupported verdicts do not disqualify. `USABLE` is reference-dependent through that conjunct, and inherits the error characterised in Section 5.1.

Sampling. With n samples per task of which c satisfy a predicate, we use the unbiased estimator of Chen et al. [2021]:

$$\widehat{\text{metric@}k} = \mathbb{E}_T \left[1 - \frac{\binom{n-c}{k}}{\binom{n}{k}} \right].$$

We draw $n = 3$ and report $k = 3$. We note explicitly that at $n = k$ this estimator degenerates to $\mathbf{1}[c \geq 1]$ and provides no variance reduction over the naive any-of- n statistic; it is reported as @3 for comparability, not because the estimator is doing work at this setting. $n > k$ would be required for the estimator’s variance properties to apply, and is left to future work with a larger budget.

3.5 Sampling configuration

Reasoning disabled. Four of eleven models emit no hidden reasoning tokens; seven do. Evaluating them together with reasoning enabled would confound model identity with inference-time compute budget. All requests therefore disable reasoning, and reasoning-token counts are recorded per response to verify compliance (observed: zero throughout).

Temperature unset. We set `require_parameters=true`, which restricts routing to providers that honour submitted sampling parameters rather than silently discarding them. Six of eleven models reject `temperature` outright; in combination the request is unroutable. We therefore omit `temperature` entirely and each model samples at its provider default. This preserves the non-silent-discard guarantee at the cost of controlled sampling—a trade we consider correct but note as a limitation.

Since $k > 1$ is only informative under output diversity, we verified diversity directly: repeated queries on a fixed task yielded 3, 3, 3 and 2 distinct patterns for `claude-opus-5`, `gpt-5.6-sol`, `kimi-k3` and `glm-5.2` respectively. The analysis pipeline additionally warns if $> 90\%$ of a model’s tasks return identical samples.

3.6 Controls

Three synthetic responses traverse the identical scoring path on every run: the task’s own reference (must score correct and usable), the pattern `z{5}` (must fail all correctness), and `(a+)+b` (must be flagged vulnerable). A run in which any control deviates is discarded rather than reported. This detects the failure mode in which a scorer silently returns null results, which is otherwise indistinguishable from uniformly poor model performance. All controls behaved as specified for all eleven models.

4 Results

4.1 Main results

The central quantitative observation in Table 1 is the gap between `PASS@3` (38.0–47.5%) and `USABLE@3` (17.1–24.8%). Approximately half of all test-passing generations fail at least one of the remaining two criteria, and the ratio is stable across the entire price range. We count 135 model-task pairs in which a pattern satisfied every test and was simultaneously screened as ReDoS-vulnerable.

`EXACT@3`—string identity with the reference—ranges from 2.0% to 6.3%, confirming that string-comparison scoring would misrank essentially all systems and justifying the automata-theoretic treatment.

Model	PASS@3	USABLE@3	VULNERABLE@3	DFA-EQ@3	DFA-EQ _{dec}	EXACT@3	undec.	fail
kimi-k3	47.2	24.8	14.1	15.2	17.5	6.3	82	16
claude-opus-5	47.5	23.0	15.5	14.0	17.0	5.4	109	36
qwen3.6-max-preview	42.4	21.6	9.1	13.8	17.4	3.8	94	0
gpt-5.6-sol	42.2	21.1	10.2	9.6	13.6	2.0	133	1
deepseek-v4-flash-0731	38.0	19.8	12.0	11.3	14.3	3.3	93	0
qwen3.6-plus	39.8	19.8	9.8	11.6	14.8	3.8	99	0
glm-5.2	42.4	18.7	14.2	10.4	12.9	3.8	86	0
gpt-5.6-luna	39.3	18.7	11.8	9.3	12.2	2.0	113	1
gpt-5.6-terra	42.2	18.7	12.0	10.2	13.1	2.0	100	0
claude-sonnet-5	40.7	18.0	10.9	10.4	13.3	3.3	97	0
gemini-3.1-flash-lite	38.7	17.1	12.0	9.3	11.8	3.1	95	0

Table 1: Primary results, 450 tasks, $n = k = 3$. All figures percentages except *undec.* (count of tasks with undecidable equivalence, of 450) and *fail* (count of failed requests, of 1,350). Ordering is by USABLE@3 and should be read as a banding, not a ranking (Section 5.3).

4.2 ReDoS and the human baseline

Table 2 applies the identical safety screen to the corpus’s 450 human-written reference patterns. For comparability with the human set (one pattern per task) we score one sample per task per model rather than the any-of-three statistic of Table 1.

Source	n	vulnerable (%)	exponential (%)	polynomial (%)
<i>Human reference answers</i>	450	13.6	6.4	7.1
kimi-k3	447	10.7	5.4	5.4
claude-opus-5	444	10.6	7.4	3.2
glm-5.2	450	10.2	5.8	4.4
gemini-3.1-flash-lite	450	9.8	5.1	4.7
gpt-5.6-sol	450	9.3	5.1	4.2
claude-sonnet-5	450	8.9	6.0	2.9
gpt-5.6-terra	450	8.7	4.9	3.8
qwen3.6-plus	450	8.4	5.1	3.3
gpt-5.6-luna	450	8.0	4.0	4.0
deepseek-v4-flash-0731	450	7.3	4.4	2.9
qwen3.6-max-preview	450	7.3	4.7	2.7
<i>All models pooled</i>	4941	9.0	5.3	3.8

Table 2: ReDoS screening of model outputs against the corpus’s own human-written reference patterns, one pattern per task in all cases.

Every model is safer than the reference set. This inverts the expected framing. The result is not that generation introduces ReDoS vulnerability relative to human practice; it is that ReDoS vulnerability is endemic to human-written regular expressions at a rate exceeding what these models reproduce.

We regard this as the paper’s most consequential empirical finding, because it changes the implied remedy. A model-deficiency framing implies waiting for better models. A distributional framing implies that safety screening must be applied at the point of use, since neither the generator nor the human corpus it was trained on performs that check.

Relation to prior work. ICPC 2024 authors — VERIFY [2024] report that LLM-generated regexes skew toward polynomial ReDoS. We do not reproduce this cleanly: pooled across models we observe 5.3% exponential against 3.8% polynomial, the opposite direction, while the human references skew polynomial (6.4% vs. 7.1%). Subcategory counts are small and we do not advance a strong claim; we note the discrepancy as warranting replication.

4.3 Cost

Model	USABLE@3 (%)	cost/task (\$ $\times 10^{-6}$)	total (\$)
deepseek-v4-flash-0731	19.8	25.6	0.03
gpt-5.6-luna	18.7	42.9	0.06
gemini-3.1-flash-lite	17.1	90.2	0.12
qwen3.6-plus	19.8	120.6	0.16
glm-5.2	18.7	158.2	0.21
qwen3.6-max-preview	21.6	387.9	0.52
gpt-5.6-terra	18.7	406.2	0.55
claude-sonnet-5	18.0	932.4	1.26
kimik-k3	24.8	1328.1	1.79
gpt-5.6-sol	21.1	2107.6	2.85
claude-opus-5	23.0	2514.2	3.39

Table 3: Cost per task at $n = 3$, measured from per-request billing returned by the serving layer rather than estimated from list prices.

The extremes differ by a factor of 98 in price and approximately three percentage points in USABLE@3—a difference at the boundary of resolvability (Section 5.2). We restrict this observation to this task; it is not a general claim about model capability.

4.4 Refusals

claude-opus-5 was blocked by a content filter on 29 requests (2.1% of its calls) across five tasks. The prompts concern: strings excluding a single quotation mark; a six-character alphanumeric password; hex byte sequences; SQL update/insert statements; and—without evident security relevance—“*Matches a file extention.*” No other model refused any request.

Refusals were *not deterministic*: several tasks were refused on one sample and answered on the other two under identical settings. This is detectable only at $k > 1$; a single-sample protocol would record a deterministic failure. We report refusals as a distinct failure category rather than as incorrect answers (Appendix A).

4.5 Inference-time compute

On the corpus’s simplest task (“*Matches exactly 1 numeric digit (0-9).*”), qwen3.6-max-preview emitted 1,571 hidden reasoning tokens before answering $\wedge[0-9]\$, at a cost of $0.0098. With reasoning disabled it produced the identical answer in 10 tokens at 1/100th the cost. We further observe that requesting *low* reasoning effort produced 2,375 reasoning tokens on that model—more than the default—indicating that effort controls are not reliably monotone.$

A reasoning-enabled comparison on 12 tasks across five models yielded per-request costs $15.7\times$ higher and no interpretable accuracy signal at that sample size (± 14 percentage points per cell). We report the cost result and explicitly decline to report an accuracy conclusion.

5 Measurement Validity

5.1 The reference standard contains errors

We first rule out a loading fault: 449 of 450 reference patterns satisfy their own tests, confirming that corpus semantics and dialect are correctly configured. The defect is subtler—a reference may satisfy its own tests and still not denote what its description states, when the tests do not exercise the discrepancy.

We drew a seeded random sample of 60 cases in which a model *passed every test* yet received a DIFF verdict, and adjudicated 14 in detail.

Adjudicated cause	count	share
Reference pattern is incorrect	5	36%
Description underspecifies the disputed property	6	43%
Model output is incorrect	3	21%

Table 4: Adjudication of 14 sampled DIFF verdicts on test-passing outputs. Per-case reasoning is released.

Independently, 19 of the 60 sampled disagreements (32%) differ only on non-ASCII input, arising from Python’s Unicode-aware `\d` versus an explicit `[0-9]`; these are formally distinct languages and practically identical. Combining both effects, we estimate the model is clearly at fault in approximately **15%** of DIFF verdicts.

Representative cases:

- **Description:** “*It just accepts only positive numbers.*” The reference `^\d+([. ,]?\d+)?$` admits 0. The model excluded zero and was scored incorrect.
- **Description:** “*A very simple ISBN validation expression—it just checks for a 10 digit number.*” The reference is `^\d{9}[\d|X]$`. The character class contains digit, *pipe* and *X*: an alternation operator written inside a class, where it denotes a literal. The reference therefore accepts `000000000|`, which is the witness separating it from the model’s `^\d{9}[\dX]$`.

- **Description:** “*...make sure commas are in the rite place (if present).*” The reference

`^\$?\d{1,3}(,?\d{3})*(\.\d{1,2})?$`

makes the comma optional, admitting `0,000000` and negating the stated purpose. The model enforced comma placement and was scored incorrect.

- **Underspecification:** “*Matches any single upper- or lower-case letter.*” The reference `^[a-zA-Z]$` requires a whole-string match; a candidate `[A-Za-z]` requires only occurrence. The description does not distinguish these. We tested whether anchoring alone explains failures generally: it accounts for approximately 6%, so the effect is diffuse rather than a single correctable convention.

Consequence. DFA-EQ and USABLE are lower bounds on model correctness by an unmodelled margin. PASS and VULNERABLE are unaffected, since neither consults the reference.

Limitation of this analysis. The adjudication comprises 14 cases judged by the present authors on their own benchmark. We release per-case reasoning with unadjudicated cases retaining an empty verdict field, and recommend that the 15% figure be re-derived from a larger sample by independent annotators before being cited quantitatively. We report it because its *direction* is unambiguous and materially affects how the composite metric should be read.

5.2 Paired inference

Our initial analysis compared systems using independent binomial confidence intervals, which is the wrong estimator for this design. All eleven models answer identical items; task difficulty is a shared random effect, and treating systems as independent samples charges each separately for it.

Under a paired bootstrap—resampling tasks with replacement, scoring all models on each resample, and forming per-pair differences—task difficulty cancels and only genuine disagreement contributes. With $B = 10,000$ resamples and a fixed seed:

Procedure	pairs resolved at 95% (of 55)
Independent binomial intervals	1
Paired bootstrap, <code>USABLE@3</code>	9
Paired bootstrap, <code>PASS@3</code>	15
Paired bootstrap, <code>VULNERABLE@3</code>	15

Table 5: Resolvable pairwise comparisons under independent versus paired inference on identical data.

The correction is not marginal: it changes the number of defensible comparisons by nearly an order of magnitude. We note this because independent intervals appear to be the default in adjacent evaluation literature, and the design—all systems, all items—is nearly universal.

Even corrected, the middle of the table does not separate. For example `qwen3.6-max-preview` versus `gpt-5.6-sol` yields +0.7% with a 95% interval of $[-2.9\%, +4.3\%]$.

5.3 Discriminative power of the corpus

A structural constraint underlies the above. On `USABLE@3`, **62% of tasks produce identical outcomes across all eleven models**; on `PASS@3`, 56%; on `VULNERABLE@3`, 77%. Approximately 167 of 450 tasks contribute any discriminative signal.

For calibration: resolving a 3-percentage-point difference at $p \approx 0.2$ with independent intervals requires roughly 2,800 items—more than the full 762-task corpus contains. Paired inference reduces this requirement substantially but cannot manufacture signal from items on which every system agrees. Benchmark construction for system comparison should therefore optimise for discriminative yield, not corpus size.

6 Threats to Validity

Contamination. The corpus predates every evaluated model and is publicly available; performance may partly reflect memorisation. We have no held-out private task set and cannot bound this. We regard it as the most significant unaddressed threat.

Single prompt. All results are conditional on one unoptimised prompt. Given the underspecification documented in Section 5.1, sensitivity to phrasing—particularly regarding match-anchoring convention—is plausibly large and is unmeasured.

Sampling control. Temperature is unset (Section 3.4), so sampling temperature varies across providers by default. Diversity was verified but not equalised.

Estimator degeneracy. At $n = k = 3$ the PASS@ k estimator reduces to any-of-3.

Coverage asymmetry. Nine models cover all 450 tasks; `kimi-k3` covers 447 (budget exhaustion mid-run) and `claude-opus-5` 445 (refusals). Denominators differ by $< 1\%$.

Screening, not proof. VULNERABLE is a lower bound (Section 3.4).

Single corpus. We do not test whether the observed banding survives corpus change.

Self-adjudication. See Section 5.1.

7 Recommendations for Evaluation Design

1. **Partition metrics by reference dependence.** Metrics evaluated by execution against inputs (PASS, VULNERABLE) are robust to reference-set defects. Metrics evaluated against a human answer key inherit its error rate. These should not be reported with equal confidence, and composites spanning both should state which conjunct carries the dependency.
2. **Sample and adjudicate disagreements.** The cost is hours. In our case it changed the paper’s conclusion.
3. **Use paired inference when all systems answer all items.** Table 5.
4. **Report discriminative yield,** not only corpus size (Section 5.3).
5. **Pin the serving provider and verify post hoc.** Record what served each request, not only what was requested (Section 3.3).
6. **Instrument the scorer with controls** that must pass and must fail (Section 3.4).
7. **Release raw generations.** This permits scoring decisions to be overturned without repurchasing inference, which is the difference between a reproducible result and an unfalsifiable one.

8 Conclusion

Across eleven contemporary models, roughly half of all regular expressions that satisfy their tests fail an additional correctness or safety criterion, and the proportion is insensitive to a $98\times$ variation in inference cost. Models are uniformly safer than the human-written reference patterns against which they are scored, which locates ReDoS prevalence in the training distribution rather than in generation.

Our second contribution is negative and, we suspect, more broadly applicable: the semantic-equivalence metric standard in this task family is, in our sample, dominated by reference-set error and task underspecification rather than by model behaviour. We were able to establish this only by auditing our own results after collecting them, and it changed what we were willing to publish. We recommend the practice generally.

Acknowledgements

Scoring is performed by `regexbench` [Foothills Labs, 2026], prior work by the present authors, used here as a pinned dependency and not contributed by this paper. The corpus is Re(gEx|DoS)Eval [S2E Lab — VERIFY, 2024], used under its original terms and not redistributed.

References

- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. Source of the unbiased pass@k estimator used in Section 3.4.
- James C. Davis, Christy A. Coghlan, Francisco Servant, and Dongyoon Lee. The impact of regular expression denial of service (redos) in practice: An empirical study at the ecosystem scale. In *Proceedings of ESEC/FSE*, 2018.
- James C. Davis, Louis G. Michael IV, Christy A. Coghlan, Francisco Servant, and Dongyoon Lee. Why aren’t regular expressions a lingua franca? an empirical study on the re-use and portability of regular expressions. In *Proceedings of ESEC/FSE*, 2019. Artifact: <https://github.com/VTLeeLab/LinguaFranca-FSE19>.
- Foothills Labs. `regexbench`: Evaluate generated regular expressions — semantic equivalence, correctness, and redos safety. <https://github.com/foothills-labs/regexbench>, 2026. Version 0.4.0, pinned at commit 05d7547b. Apache-2.0. Prior work by the present authors; used here as a dependency, not contributed by this paper.
- ICPC 2024 authors — VERIFY. On the security of regular expressions generated by large language models. In *Proceedings of the 32nd IEEE/ACM International Conference on Program Comprehension (ICPC)*, 2024. doi: 10.1145/3643916.3644424. Taxonomy of five ReDoS vulnerability families in LLM-generated patterns; reports a skew toward polynomial ReDoS.
- Nate Kushman and Regina Barzilay. Using semantic unification to generate regular expressions from natural language. In *Proceedings of NAACL-HLT*, 2013. The KB13 corpus.
- LessWrong. Not pinning your openrouter provider might invalidate your research. <https://www.lesswrong.com/posts/KsyoSAYBRXtwzSugg/>, 2025. Reports that 31 of 32 surveyed AI-safety codebases using OpenRouter did so without pinning the serving provider.
- Nicholas Locascio, Karthik Narasimhan, Eduardo DeLeon, Nate Kushman, and Regina Barzilay. Neural generation of regular expressions from natural language with minimal domain knowledge. In *Proceedings of EMNLP*, 2016. Introduces NL-RX; establishes DFA-equivalence as the standard correctness criterion for this task.

S2E Lab — VERIFY. Re(gex|dos)eval: A benchmark for regular expression generation and redos detection. <https://github.com/s2e-lab/RegexEval>, 2024. 762 natural-language-to-regex tasks with tests and reference patterns. Corpus used in this work.

A Failure taxonomy

Of 14,850 requests, 54 (0.36%) failed and are reported rather than excluded.

Model	Cause	Count
claude-opus-5	content-filter refusal	29
kimi-k3	account spending limit reached mid-run	11
kimi-k3	response without resolved provider	3
gpt-5.6-sol	response without resolved provider	1

Table 6: Request failures by cause. Responses lacking a resolved provider are discarded rather than scored, as provenance is a precondition for reproducibility.

Empty completions are classified by reported `finish_reason` into refusal, token-budget exhaustion, and unexplained, rather than being scored as empty patterns—which would be indistinguishable from a maximally poor answer.

B Operational hazards

We document configuration hazards encountered during this study, each of which produced plausible-looking but invalid data. We believe this list has independent value.

1. **Parameter combinations that are silently unroutable.** Setting `temperature` together with `require_parameters=true` yields a hard routing failure on 6 of 11 models. Our first pilot failed all 396 calls.
2. **Systems that cannot satisfy the experimental condition.** Two Gemini variants return *“Reasoning is mandatory for this endpoint”* and cannot be evaluated with reasoning disabled. Undetected, this would have placed one reasoning-enabled system in an otherwise reasoning-disabled comparison.
3. **Empty completions from budget exhaustion.** At a 200-token cap, reasoning models consumed the entire budget on hidden reasoning and returned no content.
4. **Resumption across configuration changes.** Our collection is resumable by (model, task, sample). An aborted run at one token cap left truncated rows that a resumed run at a different cap treated as complete, interleaving two configurations indistinguishably. Rows now carry a configuration fingerprint and resumption refuses on mismatch.
5. **Sample collapse in scoring.** An early scorer retained only the final sample per task, which would have rendered all @3 statistics disguised @1 statistics with no visible symptom.

C Reproduction

```
git clone https://github.com/foothills-labs/regexleaderboard
cd regexleaderboard
make setup # pinned scorer + corpus download
make score RUN=sweep
```

Scoring reads only committed generations; no API access or expenditure is required. We verified reproduction from a clean checkout with fresh dependency installation and corpus download, obtaining bit-identical metrics. A reduced form of this check executes in continuous integration on every commit.